

# An Introduction to Explainable ML

**BOSTON UNIVERSITY  
MACHINE INTELLIGENCE  
COMMUNITY**

Devin de Hueck, Jianqi Ma  
Spring, 2019

# What we will cover

- Motivation for Interpretability
- Overview of Interpretable ML Methods
- Counterfactual Explanations
- A method to generate saliency maps for black-box models
- Future Work



# The problem of explainable ML

**Why does a ML model make the prediction that it does?**





# Why do we want interpretable models?

# What is your problem?

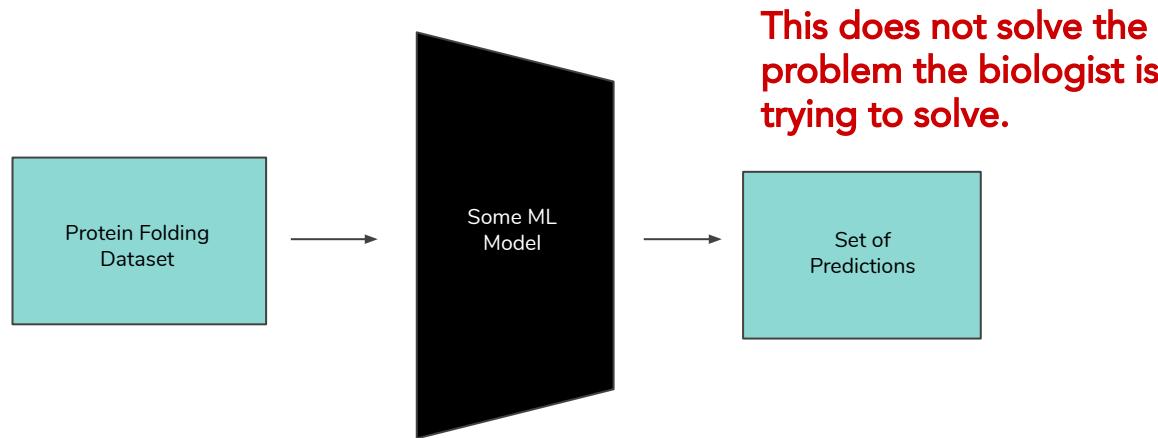
**Do you want to just know what was predicted or do you want to know why it was predicted?**



# Problems that requires interpretability

Consider the case of a biologist studying protein behavior

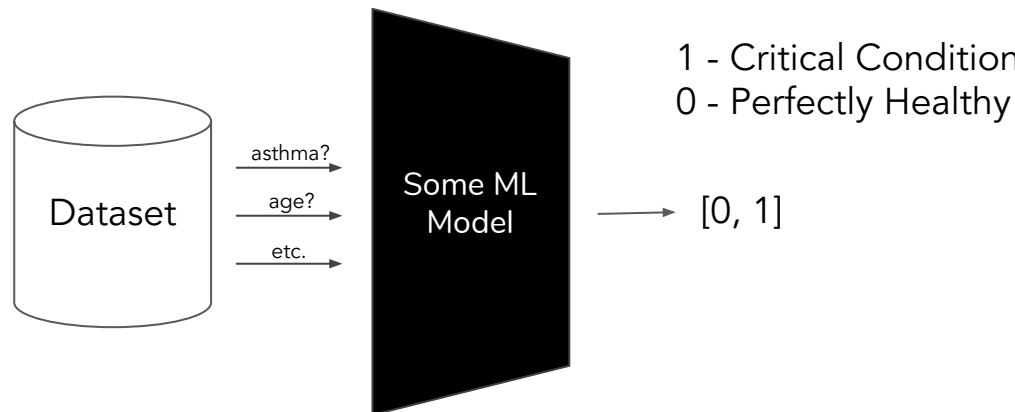
This biologist has the goal of understanding [why/how](#) proteins behave in the way they do



# Problems that requires interpretability cont.

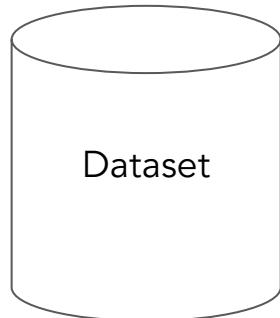
A hospital wants to predict the likelihood a patient has pneumonia.

It trains a model to predict a score between 0 and 1



# Problems that requires interpretability cont.

A hospital wants to predict the likelihood a patient has pneumonia.



Consider the case that the dataset showed that patients with asthma actually **had a better prognosis than those without**.

Imagine if the hospital was prioritizing treatment with this model. Yikes!



# Overview of Explainable ML Methods

# 3 Categories

Surrogate Models

Feature Importance

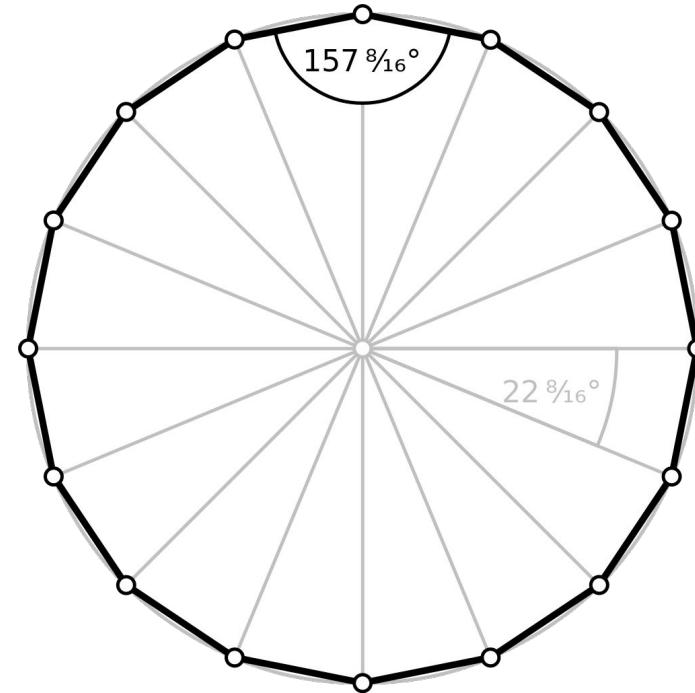
Feature Effects



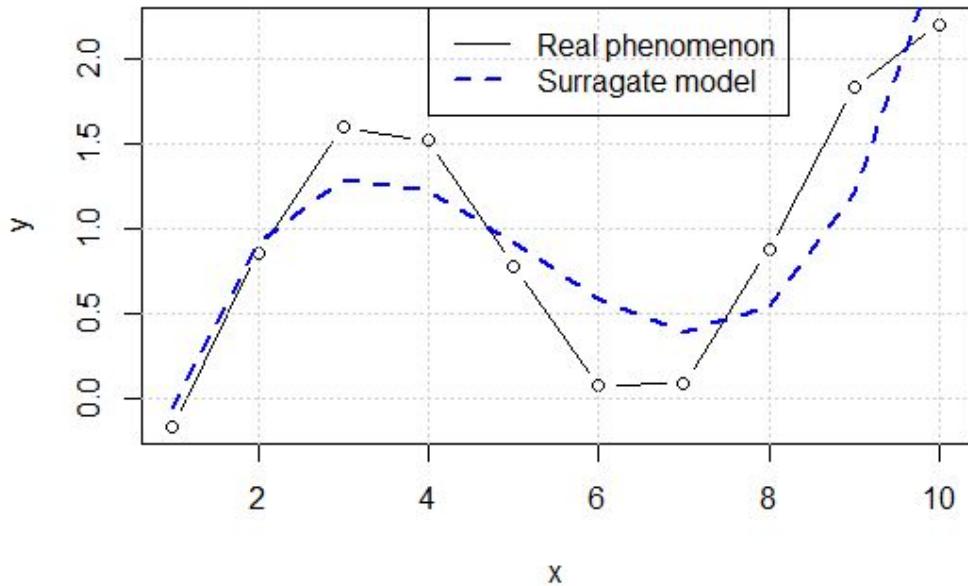
# Surrogate Models

Why using Surrogate Models:

If a model is expensive, time-consuming or otherwise difficult to measure, a **cheap** and **fast** surrogate model of the outcome can be used instead.



# Surrogate Models



The purpose of surrogate models is to approximate the predictions of the ML model as **accurately** as possible and to be **interpretable** at the same time.

# Feature Importance

Feature Importance = Increase in Prediction Error

- A feature is “important” if shuffling its values increases the model error, because in this case the model relied on the feature for the prediction.
- A feature is “unimportant” if shuffling its values leaves the model error unchanged, because in this case the model ignored the feature for the prediction.



# Feature Effects

Main Idea: summarize the average **effect** a feature has on the **prediction**

- **Partial Dependence Plot (PDP):**

The method considers all instances and gives a statement about the global relationship of a feature with the predicted outcome.

- **Individual Conditional Expectation (ICE):**

ICE plot visualizes the dependence of the prediction on a feature for each instance separately, resulting in one line per instance.

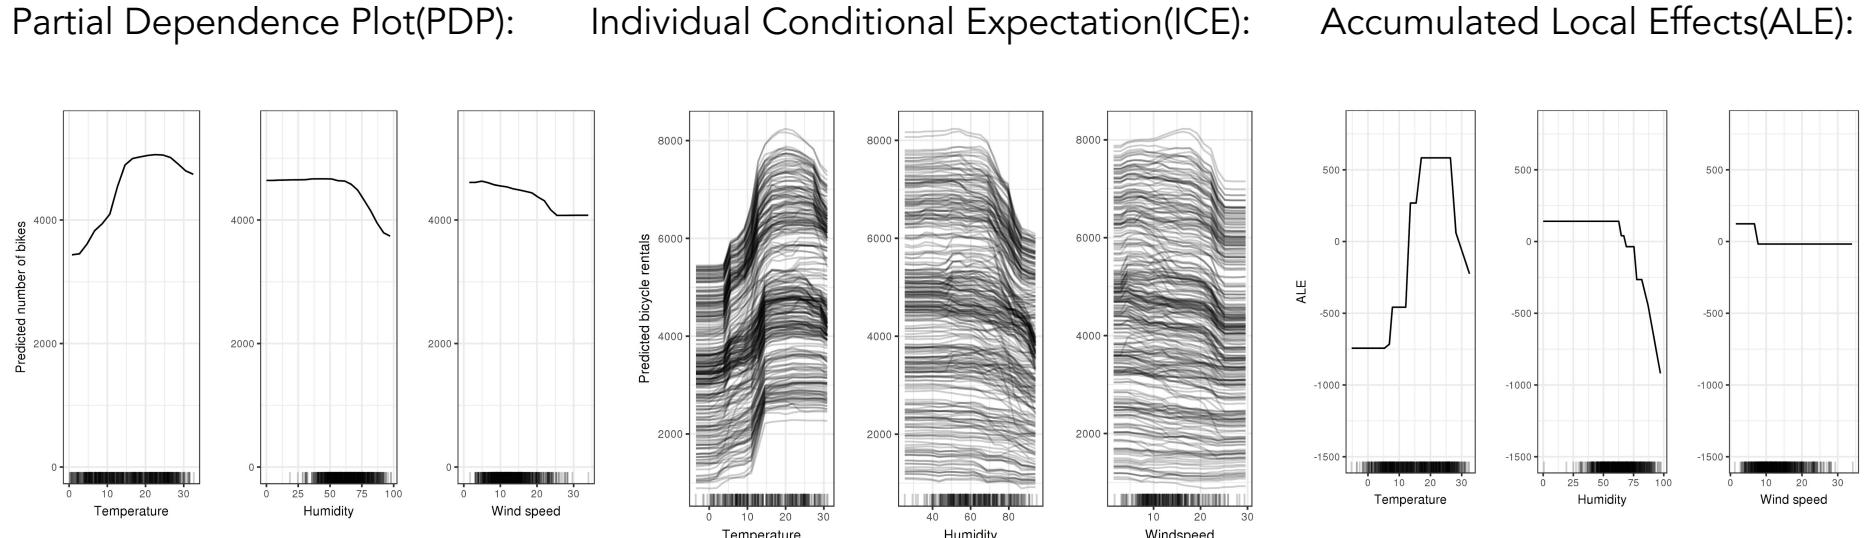
- **Accumulated Local Effects (ALE):**

Accumulated local effects describe how features influence the prediction of a machine learning model on average.

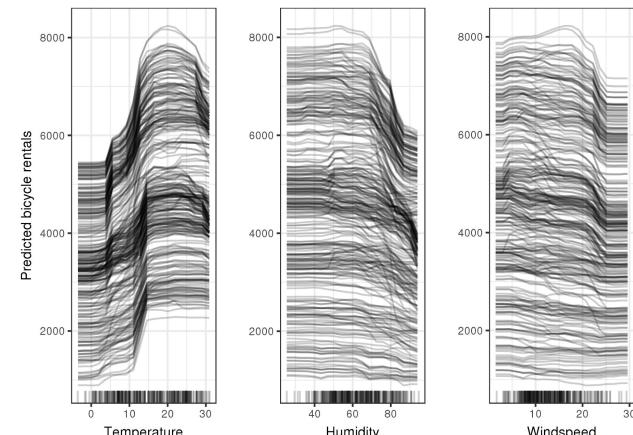


# Visualizing Feature Effects Methods

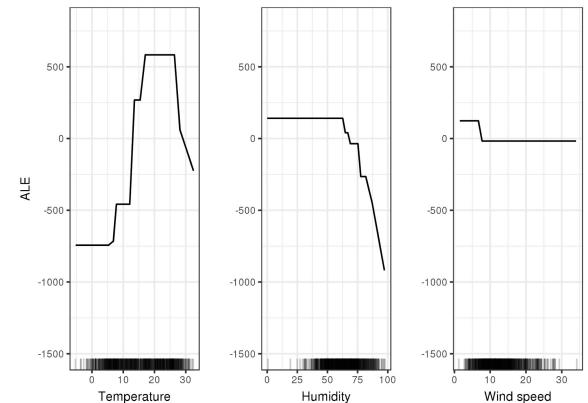
Partial Dependence Plot(PDP):



Individual Conditional Expectation(ICE):



Accumulated Local Effects(ALE):



# Properties of Explainable Models

- Global vs. Local
  - Is the interpretation for the entire model? Or maybe just for an individual prediction?
- Model-Specific vs. Model-Agnostic
  - Whitebox vs. Blackbox
- Type of Data
  - Tabular
  - Text
  - Image





# Counterfactual Explanations

# What is a Counterfactual Explanation?

**If X had not occurred, Y would not have occurred**

We contradict observed facts to “imagine” a reality in which these facts did not occur

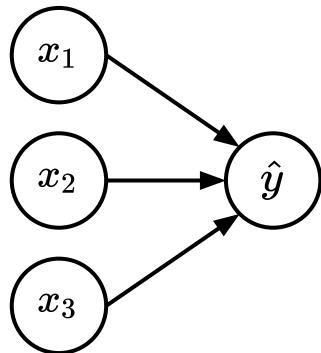
E.g. If I hadn’t taken a sip of this hot coffee, I wouldn’t have burned my tongue



# What is a Counterfactual Explanation in the Context of ML?

Our inputs **cause** our outputs  
so an input **explains** an output.

We explain one output with examples  
- Example Based/Local



# Some Scenarios for Counterfactual Explanations

**Given a model** which outputs a decision of denial or acceptance for a loan based on factors such as income, assets, age, etc.

**We ask the question:** What would have to change for an individual who has been rejected to be granted the loan?

**Plausible Counterfactual:**

"You were denied a loan because your annual income was \$30,000. If your income had been \$45,000, you would have been offered a loan."



# Some Scenarios for Counterfactual Explanations

**Given a model** which outputs a diabetes risk score for a woman of Pima heritage based on factors such as age, BMI, insulin level, etc.

**We ask the question:** What would have to change for an individual to have a risk score of 0.5?

**Plausible Counterfactual:**

"If the individual's 2-Hour serum insulin level was 169.5, the risk score would be 0.51"

*This is already very similar to how doctors communicate risks. E.g. "A BMI greater than X indicates obesity and an increase in health risks."*



# Generating Counterfactuals - Notation

**We asked the question:** *What would have to change for an individual to have a risk score of 0.5?*

**General Case:** What in the input  $x$  would have to change such that the output has the desired outcome  $y'$ ?

$x'$  - The counterfactual that we are generating

$x$  - An input value

$y'$  - The desired outcome

$x'$  is the smallest change in  $x$  such that the prediction is the desired outcome  $y'$ .

# Generating Counterfactuals - The Loss Function

Our assertion that  $x'$  is the smallest change in  $x$  such that the prediction is the desired outcome  $y'$  is characterized by the loss function:

$$L(x, x', y', \lambda)$$

# Generating Counterfactuals - The Loss Function

Our assertion that  $x'$  is the smallest change in  $x$  such that the prediction is the desired outcome  $y'$  is characterized by the loss function:

$$L(x, x', y', \lambda) = \lambda \cdot (\hat{f}(x') - y')^2$$

The hyperparameter  
lambda

Distance between the  
prediction from the  
current counterfactual  
and desired outcome



# Generating Counterfactuals - The Loss Function

Our assertion that  $x'$  is the smallest change in  $x$  such that the prediction is the desired outcome  $y'$  is characterized by the loss function:

$$L(x, x', y', \lambda) = \frac{\lambda \cdot (\hat{f}(x') - y')^2}{\text{Distance between the prediction from the current counterfactual and desired outcome}} + \frac{d(x, x')}{\text{Distance between the counterfactual example and the original example}}$$

The diagram illustrates the components of the loss function. The first term,  $\lambda \cdot (\hat{f}(x') - y')^2$ , is divided by a horizontal line with an upward-pointing arrow. This arrow points to the text "Distance between the prediction from the current counterfactual and desired outcome". The second term,  $d(x, x')$ , is also divided by a horizontal line with an upward-pointing arrow. This arrow points to the text "Distance between the counterfactual example and the original example". The hyperparameter  $\lambda$  is indicated by an arrow pointing to its position in the term  $\lambda \cdot (\hat{f}(x') - y')^2$ .

# Generating Counterfactuals - The Loss Function

Our assertion that  $x'$  is the smallest change in  $x$  such that the prediction is the desired outcome  $y'$  is characterized by the loss function:

$$L(x, x', y', \lambda) = \lambda \cdot (\hat{f}(x') - y')^2 + d(x, x')$$

↑  
The hyperparameter  
lambda

↑  
Distance between the  
prediction from the  
current counterfactual  
and desired outcome

↑  
Distance between the  
counterfactual  
example and the  
original example

Find the values of  $x'$  and  $\lambda$  such that:  $\arg \min_{x'} L(x, x', y', \lambda)$

# Generating Counterfactuals - The Algorithm

## 1. Initialize:

$x$  - You input that is to be explained

$y'$  - The desired outcome

$\epsilon$  - The tolerance (an acceptable distance for  $f(x')$  to be from  $y'$ )

$\lambda$  - Hyperparameter to balance the two parts of the loss function (usually set to a low value,  $> 1$ )

$x'$  - A random input instance



# Generating Counterfactuals - The Algorithm

2. Minimize the loss function (find an  $x'$  really close to  $x$ ):

$$L(x, x', y', \lambda) = \lambda \cdot (\hat{f}(x') - y')^2 + d(x, x')$$



# Generating Counterfactuals - The Algorithm

2. Minimize the loss function (find an  $x'$  really close to  $x$ ):

$$L(x, x', y', \lambda) = \lambda \cdot (\hat{f}(x') - y')^2 + d(x, x')$$

Until satisfied:

Use an optimization algorithm like gradient descent to update  $x'$  to minimize the loss:

$$x' = x' - \alpha \nabla_{x'} L(x, x', y', \lambda)$$

# Generating Counterfactuals - The Algorithm

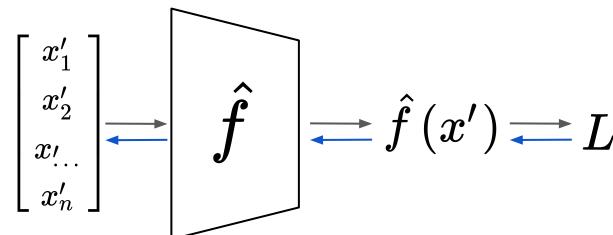
2. Minimize the loss function (find an  $x'$  really close to  $x$ ):

$$L(x, x', y', \lambda) = \lambda \cdot (\hat{f}(x') - y')^2 + d(x, x')$$

Until satisfied:

Use an optimization algorithm like gradient descent to update  $x'$  to minimize the loss:

$$x' = x' - \alpha \nabla_{x'} L(x, x', y', \lambda)$$



# Generating Counterfactuals - The Algorithm

3. Maximize  $\lambda$  until we have our desired outcome while maintaining a  $x'$  similar to  $x$

**While**  $|f(x') - y'| > \epsilon$

$\lambda = \lambda + \alpha$  - Increase lambda by a small step size  $\alpha$



# Generating Counterfactuals - The Algorithm

3. Maximize  $\lambda$  until we have our desired outcome while maintaining a  $x'$  similar to  $x$

While  $|f(x') - y'| > \epsilon$

$\lambda = \lambda + \alpha$  - Increase lambda by a small step size  $\alpha$

Use an optimization algorithm like gradient descent to update  $x'$  to minimize the loss:

$$x' = x' - \alpha \nabla_{x'} L(x, x', y', \lambda)$$



# Generating Counterfactuals - The Algorithm

3. Maximize  $\lambda$  until we have our desired outcome while maintaining a  $x'$  similar to  $x$

**While**  $|f(x') - y'| > \epsilon$

$\lambda = \lambda + \alpha$  - Increase lambda by a small step size  $\alpha$

Use an optimization algorithm like gradient descent to update  $x'$  to minimize the loss:

$$x' = x' - \alpha \nabla_{x'} L(x, x', y', \lambda)$$

**Return  $x'$  and Repeat**



# What is a Good Counterfactual Explanation?

$x'$  is the smallest change in  $x$  such that the prediction is the desired outcome  $y'$ .

All features must be of reasonable scale

- The counterfactual: "If the apartment had 300 bathrooms the rent prediction would 1500/month" is not a meaningful counterfactual

Limit features to be within the max and min values in the dataset



# Advantages

- Very clear explanations
  - Humans understand these type of explanations very well
- Does not require much domain knowledge to understand
  - Don't need to be a doctor to understand a counterfactual concerning a medical model

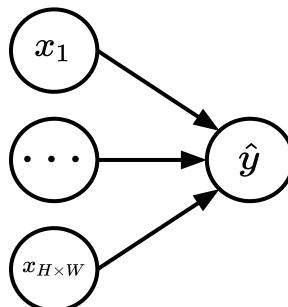
# Disadvantages

- Rashmon Effect
  - There can be multiple *valid* counterfactual examples for each input
- No guarantee of finding a counterfactual within the set tolerance ( $\epsilon$ )

# Adversarial Examples vs Counterfactual Explanations

$x'$  is the smallest change in  $x$  such that the prediction is the desired outcome  $y'$ .

What happens when our input is an image? Finding a counterfactual would give us an adversarial example.



“panda”  
57.7% confidence

<https://openai.com/blog/adversarial-example-research/>

$$\text{“panda”} + \epsilon = \text{“gibbon”}$$

“gibbon”  
99.3% confidence

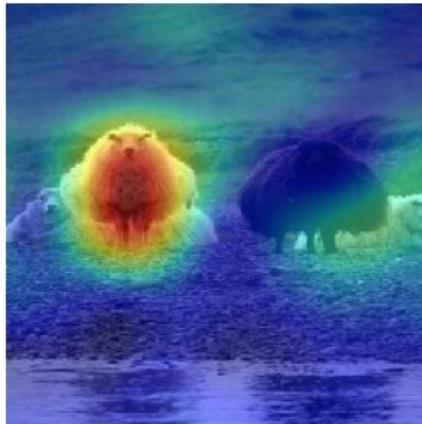


# Saliency Map Generation

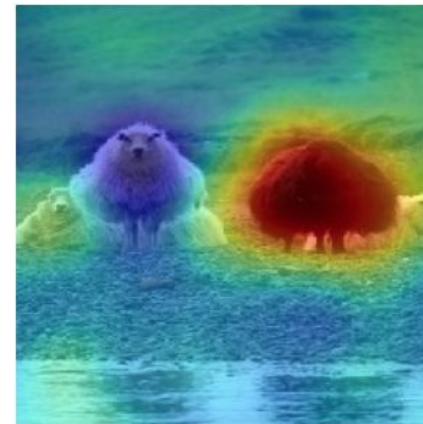
# What Are Saliency Maps?

Saliency maps tell us how important each pixel of an image is for a prediction

Saliency map for prediction of class **sheep**

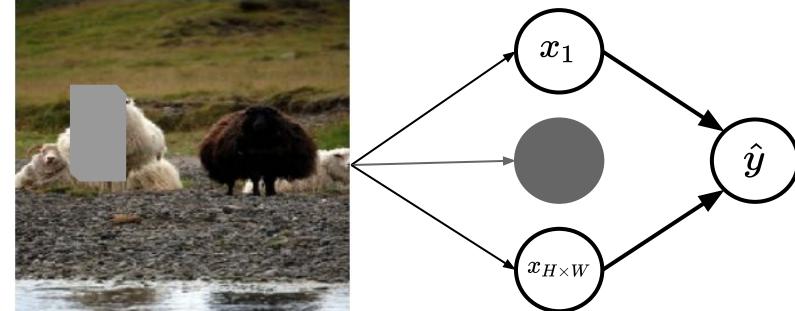
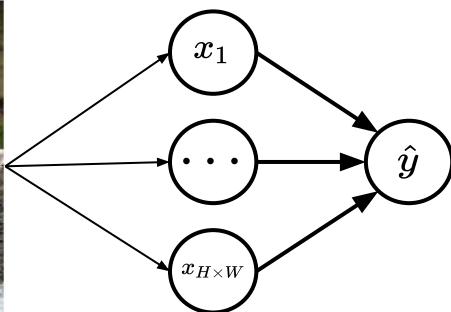


Saliency map for prediction of class **cow**



# Causality and Saliency Maps

What will happen if we delete some inputs/pixels? How will the predicted class score change?

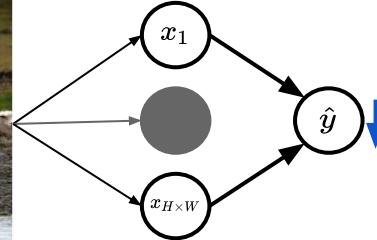
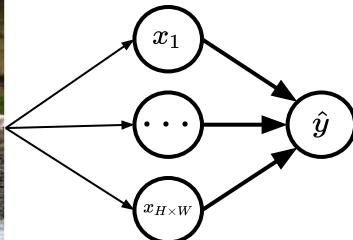


# Causality and Saliency Maps

What will happen if we delete some inputs/pixels? How will the predicted class score change? **It depends!**

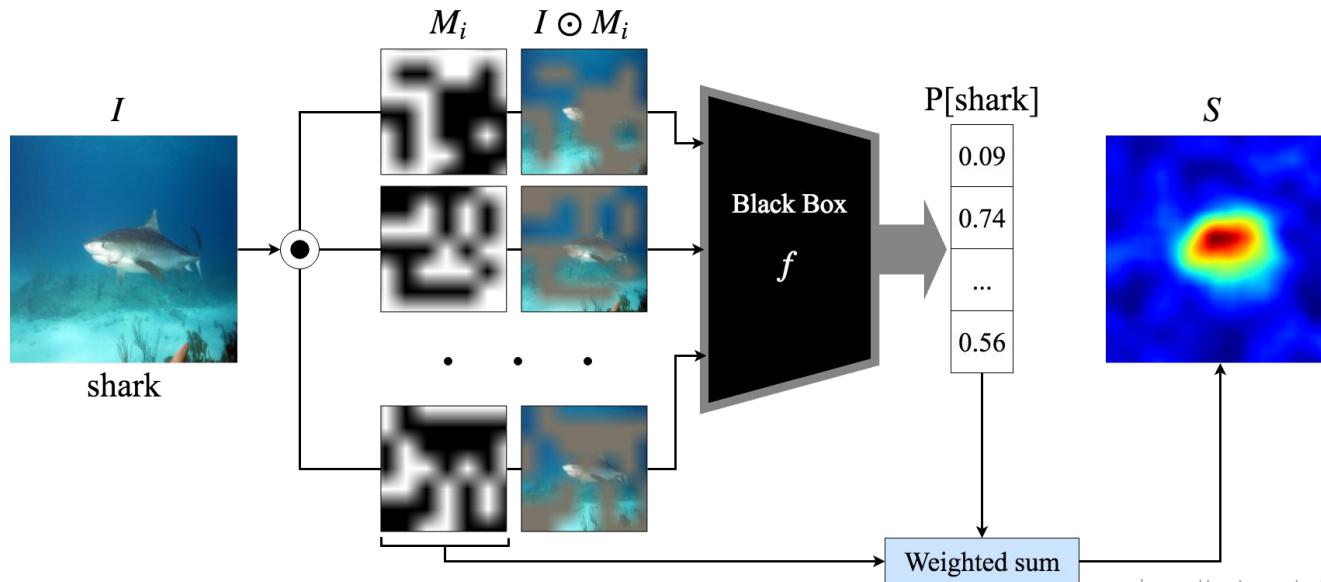
If the inputs/pixel that were removed **were important** for making the prediction then the predicted class score would **decrease** based on the importance.

If the inputs/pixel that were removed **were not important** for making the prediction then the predicted class score would **stay about the same**.



# RISE: Randomized Input Sampling for Explanation of Black-box Models

If we randomly remove pixels from different parts of an image and make a prediction on these modified images then we will eventually learn the importance of each pixel.



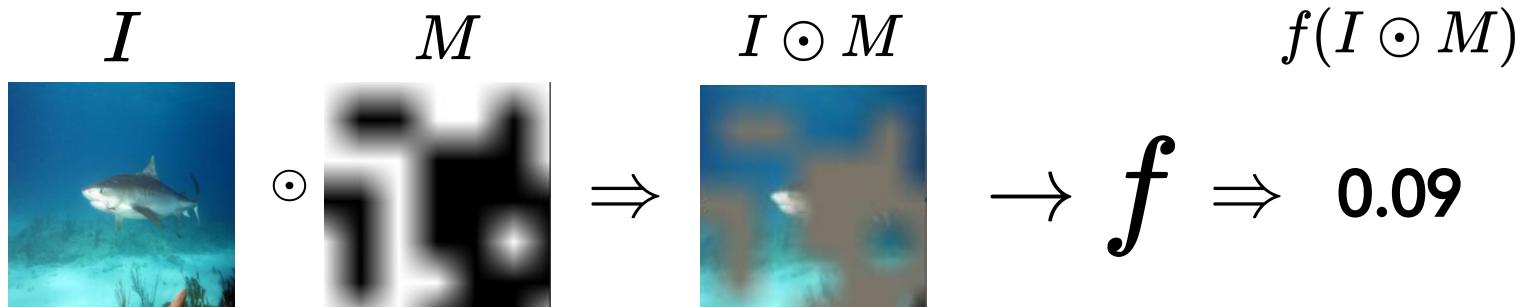
# RISE - Notation

$I$  - The image to be explained which is of size  $H \times W \times 3$

$M$  - A binary mask which is of size  $H \times W$

$I \odot M$  - A masked image which is of size  $H \times W \times 3$

$f(I \odot M)$  - The prediction from a masked image



# RISE - The importance of a pixel

"We define importance of pixel  $\lambda$  as the expected score over all possible masks  $M$  conditioned on the event that pixel  $\lambda$  is observed, i.e.,  $M(\lambda) = 1$ "

$$S_{I,f}(\lambda)$$



Saliency/Importance  
of pixel  $\lambda$

# RISE - The importance of a pixel

"We define importance of pixel  $\lambda$  as the expected score over all possible masks  $M$  conditioned on the event that pixel  $\lambda$  is observed, i.e.,  $M(\lambda) = 1$ "

$$\underline{S_{I,f}(\lambda)} = \mathbb{E}_M[\underline{f(I \odot M)}]$$

↑  
Saliency/Importance  
of pixel  $\lambda$

↑  
Black-box model's  
prediction of a  
masked image

# RISE - The importance of a pixel

"We define importance of pixel  $\lambda$  as the expected score over all possible masks  $M$  conditioned on the event that pixel  $\lambda$  is observed, i.e.,  $M(\lambda) = 1$ "

$$S_{I,f}(\lambda) = \mathbb{E}_M[f(I \odot M) | M(\lambda) = 1]$$

↑  
Saliency/Importance  
of pixel  $\lambda$

↑  
Black-box model's  
prediction of a  
masked image

↑  
Visibility of pixel  $\lambda$  - 0  
if hidden by mask, 1 if  
visible

# RISE - The importance of a pixel

"We define importance of pixel  $\lambda$  as the expected score over all possible masks  $M$  conditioned on the event that pixel  $\lambda$  is observed, i.e.,  $M(\lambda) = 1$ "

$$S_{I,f}(\lambda) = \mathbb{E}_M[f(I \odot M) | M(\lambda) = 1]$$

↑  
Saliency/Importance  
of pixel  $\lambda$

↑  
Black-box model's  
prediction of a  
masked image

↑  
Visibility of pixel  $\lambda$  - 0  
if hidden by mask, 1 if  
visible

***To get the complete saliency map we solve this equation for every pixel in the image***

# RISE - The importance of a pixel - In practice

In practice we perform a Monte-Carlo simulation.

The importance of one pixel is the average predicted score of N masked images where that pixel is visible normalized by the expected value of a mask (0.5).

$$S_{I,f}(\lambda) \stackrel{\text{MC}}{\approx}$$



# RISE - The importance of a pixel - In practice

In practice we perform a Monte-Carlo simulation.

The importance of one pixel is the average predicted score of N masked images where that pixel is visible normalized by the expected value of a mask (0.5).

$$S_{I,f}(\lambda) \stackrel{\text{MC}}{\approx} \sum_i f(I \odot M_i)$$

---

Sum of masked images  
prediction scores

# RISE - The importance of a pixel - In practice

In practice we perform a Monte-Carlo simulation.

The importance of one pixel is the average predicted score of N masked images where that pixel is visible normalized by the expected value of a mask (0.5).

$$S_{I,f}(\lambda) \stackrel{\text{MC}}{\approx} \sum_i f(I \odot M_i) \cdot \underline{M_i(\lambda)}$$

Only want to sum if visible



# RISE - The importance of a pixel - In practice

In practice we perform a Monte-Carlo simulation.

The importance of one pixel is the average predicted score of N masked images where that pixel is visible normalized by the expected value of a mask (0.5).

$$S_{I,f}(\lambda) \stackrel{\text{MC}}{\approx} \frac{1}{N} \sum_{i=1}^N f(I \odot M_i) \cdot M_i(\lambda)$$

---

Get the average  
predicted score



# RISE - The importance of a pixel - In practice

In practice we perform a Monte-Carlo simulation.

The importance of one pixel is the average predicted score of N masked images where that pixel is visible normalized by the expected value of a mask (0.5).

$$S_{I,f}(\lambda) \stackrel{\text{MC}}{\approx} \frac{1}{\mathbb{E}[M]} \cdot \frac{1}{N} \sum_{i=1}^N f(I \odot M_i) \cdot M_i(\lambda)$$

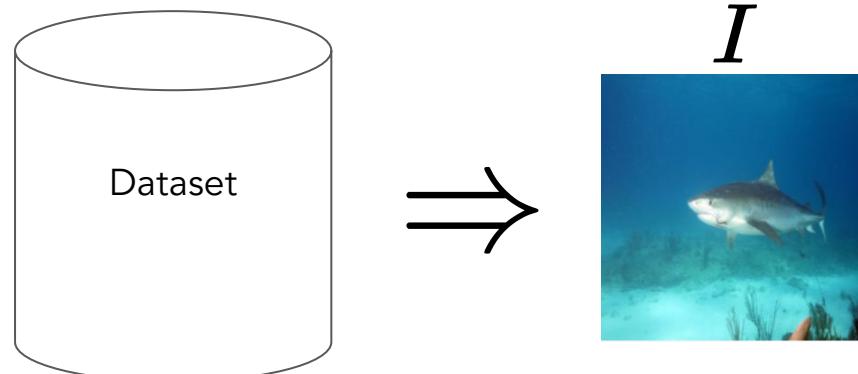
Normalize by the  
expected value of a  
mask (0.5)

The average predicted  
score of N masked  
images



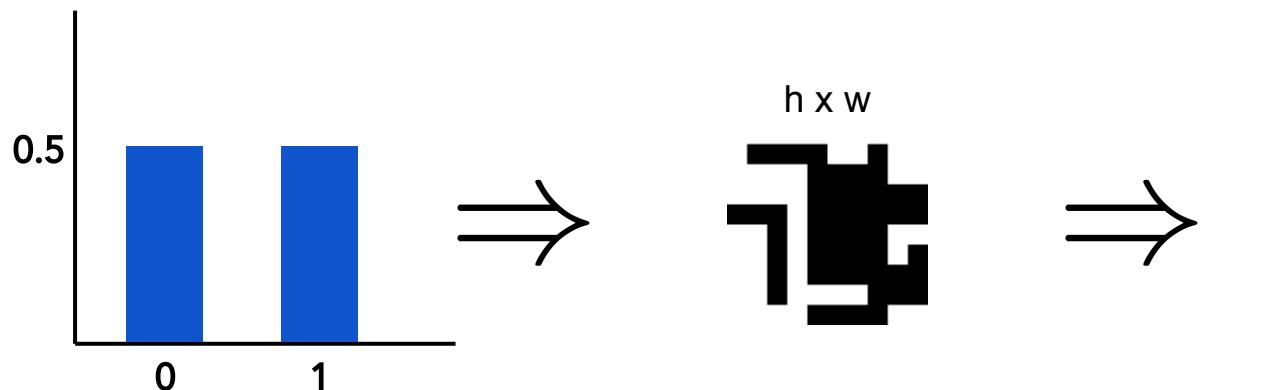
# RISE - The algorithm

## 1. Select and image to have explained



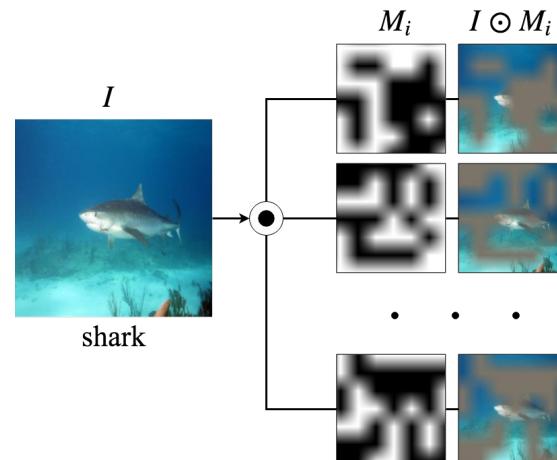
# RISE - The algorithm

## 2. Generate a set of N masks



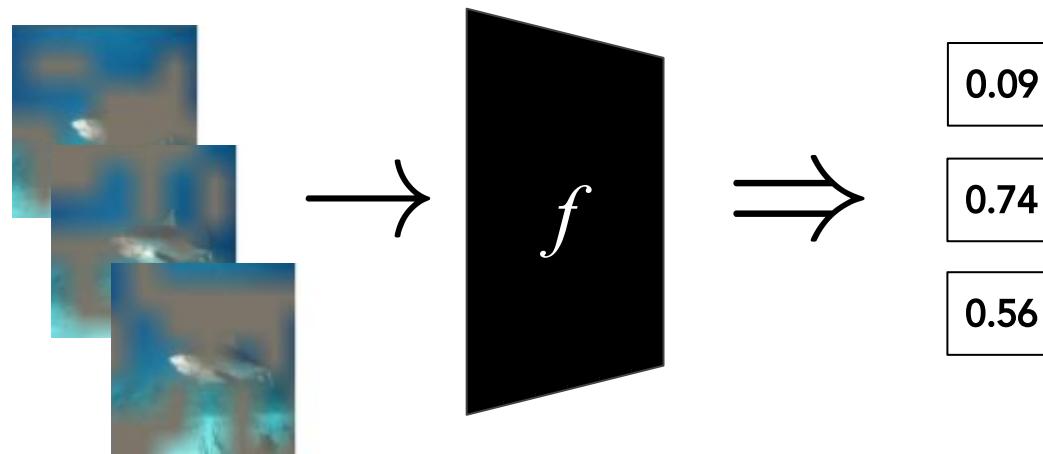
# RISE - The algorithm

3. Generate N masked images through element-wise multiplication of a mask and the image.



# RISE - The algorithm

4. Get N predictions from your black-box model on the each masked image



# RISE - The algorithm

5. Take the dot product between the N masks and N predictions then normalize to get one saliency map for each possible prediction

$$\begin{array}{c} \text{[Blurry Mask]} \cdot \boxed{0.09} + \text{[Blurry Mask]} \cdot \boxed{0.74} + \cdots + \text{[Blurry Mask]} \cdot \boxed{0.56} \\ \hline = \quad \text{[Saliency Map] } S \end{array}$$

$\mathbb{E}[M] \cdot N$

*And that's it!*

A decorative graphic on the left side of the slide features a network of nodes and lines. The nodes are small circles in shades of orange, pink, and purple, connected by thin, curved lines that form a organic, branching pattern.

# What does the future hold?

# The Future of Explainable ML

**Focus will be on Model-Agnostic methods**

**Explainable ML add-ons to currently existing libraries**



# References & Further Reading

1. <https://christophm.github.io/interpretable-ml-book/counterfactual.html>
2. Petsiuk, Vitali, Abir Das, and Kate Saenko. "[RISE: Randomized Input Sampling for Explanation of Black-box Models.](#)" arXiv preprint arXiv:1806.07421 (2018).
3. Doshi-Velez, Finale, and Been Kim. "[Towards a rigorous science of interpretable machine learning.](#)" arXiv preprint arXiv:1702.08608 (2017).
4. Hall, Patrick, and Navdeep Gill. [An Introduction to Machine Learning Interpretability](#). O'Reilly Media Inc., 2018.



# RISE - The algorithm

1. Select an image to have explained
2. Generate N masks
3. Generate N masked images through element-wise multiplication of a mask and the image.
4. Get N predictions from your black-box model on the each masked image
5. Take the dot product between the N masks and N predictions then normalize to get one saliency map for each possible prediction